ChatGPT

Frontier Research on Memory Mechanisms for Autonomous Agents

Executive Summary

From 2023 to early 2026, the center of gravity in autonomous-agent memory research shifted from simple retrieval buffers and reflective text traces toward **structured, typed, temporally aware, and increasingly learned memory systems**. Early influential work such as *Generative Agents*, *Voyager*, *Reflexion*, *MemoryBank*, and *MemGPT* established the main design patterns: natural-language episodic logs, reflection-derived summaries, procedural skill libraries, and tiered "OS-like" context management. Recent systems push further by adding graph structure, temporal semantics, multimodal memory, selective consolidation, and reinforcement-learned memory control. The strongest frontier trend is not "more memory" in the abstract; it is **better memory operations**: what to write, when to consolidate, how to preserve provenance, how to resolve contradictions, and how to retrieve action-relevant context under tight latency and token budgets. [1]

The most consequential recent papers are those that move beyond heuristic storage. *A-MEM* turns memory into a dynamically linked note network inspired by Zettelkasten; *MemoryOS* introduces hierarchical short-, mid-, and long-term stores; *Mem0* emphasizes production-scale extraction and low-latency retrieval; *G-Memory* extends structure to multi-agent collaboration; *MIRIX* introduces six typed memory classes for multimodal assistants; *Temporal Semantic Memory* models durative and time-valid memories; *LightMem* separates fast online retrieval from slower offline consolidation; *Memory-R1* learns ADD/UPDATE/DELETE/NOOP policies with RL; *APEX-MEM* uses append-only temporal property graphs; and *MemMachine* argues for preserving raw episodic ground truth and moving more intelligence into retrieval rather than lossy ingestion. [2]

Benchmarking also matured substantially. The field progressed from **LoCoMo** and **LongMemEval**, which focus on long-horizon conversational memory, to **MemoryAgentBench** for multi-turn competencies, **MemoryArena** for memory-in-the-loop action, **Mem-Gallery** for multimodal memory, **LoCoMo-Plus** for latent "cognitive memory" constraints, and **LifelongAgentBench** for continual, interdependent tasks across database, OS, and knowledge-graph environments. This is a major advance: newer benchmarks increasingly test whether memory **improves future decisions**, not merely whether an agent can quote old facts. [3]

On current evidence, the most promising research directions are: **typed memory decomposition** (episodic, semantic, procedural, profile, resource), **temporal and append-only representations** for evolving facts, **learned memory control policies**, **multimodal memory**, and **benchmarks that couple memory to action and continual adaptation**. The biggest unresolved problems are still memory write quality, temporal conflict resolution, selective forgetting, privacy/governance, cross-benchmark comparability, and reproducibility across different backbones, judges, and token budgets. [4]

I prioritized primary sources from arXiv [5] and official GitHub [6] repositories, plus official project documentation and release notes where relevant. [7]

Landscape and Taxonomy

A useful way to read the literature is to separate **memory representation** from **memory control**. Representation asks whether memory is stored as raw episodes, summaries, symbolic notes, graph nodes/edges, code skills, or multimodal artifacts. Control asks whether storage, update, deletion, consolidation, and retrieval are hand-designed heuristics, prompted by the LLM at test time, or optimized with supervised or reinforcement learning. Across 2023–2026, the field moved from “retrieve relevant chunks” toward “maintain an evolving memory substrate with lifecycle operations.” [8]

The main memory types now appearing repeatedly are:

- **Working memory**: the currently available context window or a compact in-session cache.
- **Episodic memory**: raw or lightly processed interaction traces, trial histories, trajectories, or conversation segments.
- **Semantic memory**: distilled facts, user profiles, entities, relations, project knowledge, and high-level abstractions.
- **Procedural memory**: executable skills, policies, workflows, reflections-as-rules, or action patterns.
- **Temporal memory**: memories indexed by event time, duration, supersession, or validity intervals rather than only dialogue order.
- **External memory**: vector stores, SQL/FTS indices, graphs, files, or knowledge bases outside model weights.
- **Consolidated memory**: promoted long-term knowledge formed from repeated evidence, reflection, or offline summarization. [9]

Most frontier architectures now look like this:

```
flowchart LR
    O[Observations and tool traces] --> W[Working memory]
    W --> E[Episodic store]
    E -->|summarize / distill| S[Semantic memory]
    E -->|extract skills / policies| P[Procedural memory]
    S --> R[Retriever / router]
    P --> Planner[Planner or policy]
    R --> Planner
    Planner --> A[Actions / dialogue / tool use]
    A --> F[Feedback / outcomes]
    F --> C[Write-update-delete-consolidate policy]
    C --> W
    C --> E
    C --> S
    C --> P
```

This flow is a synthesis of the dominant designs in *Generative Agents*, *Voyager*, *Reflexion*, *MemGPT*, *A-MEM*, *MemoryOS*, *MIRIX*, *LightMem*, and *Memory-R1*: a short-lived working layer, one or more persistent stores, and an explicit control loop for storage, retrieval, and consolidation. [10]

The strongest analytical distinction in the recent literature is between **heuristic memory management** and **policy-learned memory management**. Heuristic systems such as *MemGPT*, *MemoryOS*, *Mem0*, *A-MEM*, and *LightMem* encode hand-designed storage/consolidation logic, even when the LLM helps with extraction. Learned systems such as *Memory-R1* and the 2025–2026 RL line represented by *Learn to Memorize* and *Mem-α* treat memory operations as decisions that can be optimized by reward rather than manually specified. That shift is likely to matter as agent environments become more interactive, long-horizon, and multimodal. [11]

Representative Methods and Comparison Table

The table below emphasizes methods that either changed the field's design vocabulary or materially advanced state of the art. An em dash in the **Code** column means I am not citing an official public repository here, not necessarily that private code does not exist. Reported results are taken from the authors' own papers; because backbones, subsets, judges, and budgets differ, they should be read as **directional** rather than perfectly apples-to-apples. [12]

Paper	Year	Link	Memory type	Capacity / scaling	Training paradigm	Tasks	Key results
Generative Agents	2023	paper [13]	natural-language episodic stream + reflection	full experience log with dynamic retrieval	prompted memory/ reflection, no finetuning	social simulation, long-horizon planning	25-agent sandbox; coordinated party behavior emerges from one seed instruction [14]
Voyager	2023	paper [16]	procedural skill memory	ever-growing executable code library	iterative prompting, self-verification, no weight updates	embodied lifelong learning	3.3x more unique items, 2.3x farther travel, up to 15.3x faster tech milestones [17]
Reflexion	2023	paper [19]	episodic reflective memory	compact cross-trial buffer	verbal RL via feedback, no finetuning	coding, reasoning, sequential decision tasks	91% pass@1 on HumanEval vs reported 80% for GPT-4 baseline [20]

Paper	Year	Link	Memory type	Capacity / scaling	Training paradigm	Tasks	Key results
MemoryBank	2023	paper [22]	long-term semantic + episodic with forgetting	continuous updates with Ebbinghaus-inspired forgetting	prompted storage/ update heuristics	companion chat, personalization	improved empathy, relevant recall, and user-personality understanding in SiliconFriend [22]
MemGPT	2023	paper [24]	tiered working + long-term external memory	“virtual context management” across memory tiers	OS-inspired controller	document analysis, multi-session chat	handles documents beyond context window; chat agents can remember and evolve across sessions [25]
ExpeL	2023 / AAAI 2024	paper [27]	experiential memory + distilled insights + examples	grows with number of training experiences	autonomous experience gathering, no parametric updates	decision-making and transfer across training tasks	performance improves as the agent accumulates experiences; transfer emerges qualitatively [28]
A-MEM	2025	paper [30]	agentic semantic graph / linked notes	dynamic indexing, linking, and evolving notes	agent-driven organization with LLM assistance	general LLM-agent memory	beats prior baselines across six foundation models [31]
MemoryOS	2025	paper [33]	hierarchical STM / MTM / LTM	FIFO short→mid, segmented paging mid→long	heuristic OS-style management	long-term conversational personalization	+49.11% F1 and +46.18% BLEU-1 on LoCoMo over baselines on GPT-4o-mini [34]

Paper	Year	Link	Memory type	Capacity / scaling	Training paradigm	Tasks	Key results
Mem0	2025	paper [36]	extracted semantic memory; optional graph memory	production-oriented persistent store with graph variant	heuristic extraction / consolidation	conversational memory, personalization	+26% relative judge improvement over OpenAI memory system; graph variant adds ~2%; 91% lower p95 latency and >90% token savings [37]
G-Memory	2025	paper [39]	hierarchical multi-agent graph memory	three-tier insight / query / interaction graphs	evolving graph memory for MAS	multi-agent QA, embodied action	up to +20.89% embodied success and +10.12% QA accuracy across five benchmarks [40]
MIRIX	2025	paper [42]	six typed memories + multi-agent controller	Core, Episodic, Semantic, Procedural, Resource, Knowledge Vault	modular typed-memory control	multimodal personal assistant, ScreenshotVQA, LoCoMo	+35% accuracy over RAG and 99.9% storage reduction on ScreenshotVQA; 85.4% on LoCoMo [43]
Temporal Semantic Memory	2026	paper [45]	temporal semantic + durative memory	semantic timeline with duration-aware consolidation	temporal-aware retrieval and construction	personalized dialogue memory	up to 12.2% absolute accuracy improvement on LongMemEval and LoCoMo [46]
LightMem	2026	paper [47]	STM / MTM / LTM + offline consolidation	fixed online retrieval budget; supports multi-user IDs	SLM-assisted online retrieval + offline consolidation	long-horizon conversation	~2.5 average F1 gain over A-MEM on LoCoMo; 83 ms retrieval and 581 ms end-to-end latency [48]

Paper	Year	Link	Memory type	Capacity / scaling	Training paradigm	Tasks	Key results
Memory-R1	2025 / 2026	paper [50]	learned external memory ops + answer distillation	ADD / UPDATE / DELETE / NOOP; scales 3B–14B	RL with PPO / GRPO using 152 QA pairs	dialogue QA, transfer to MSC and LongMemEval	+28% F1, +34% BLEU-1, +30% judge over Mem0 on LoCoMo with LLaMA-3.1-8B-Instruct [51]
APEX-MEM	2026	paper [52]	append-only temporal property graph	full-history preservation with retrieval-time conflict resolution	multi-tool retrieval agent	long-term conversational AI	88.88% LoCoMo QA and 86.2% LongMemEval [53]
MemMachine	2026	paper [54]	ground-truth-preserving episodic + profile + short-term memory	stores whole conversational episodes; retrieval expands nucleus matches with local context	retrieval-centric, low-loss ingestion	personalized assistants, long-term QA	0.9169 on LoCoMo, 93.0% on LongMemEval-S, ~80% fewer input tokens than Mem0 under matched conditions [55]

The code ecosystem is now strong enough that reproducible systems research is plausible, but it is unevenly distributed. Official public repositories are readily available for *Voyager*, *Reflexion*, *MemoryBank*, *MemGPT*, *ExpeL*, *A-MEM*, *MemoryOS*, *Mem0*, *G-Memory*, *MIRIX*, *LightMem*, and the benchmark suites themselves. By contrast, some of the strongest recent preprints I reviewed, including *Temporal Semantic Memory*, *Memory-R1*, and *APEX-MEM*, did not have an official public repository cited in the primary materials surfaced here. [57]

A separate but important production layer has also emerged around persistent-agent infrastructure rather than benchmark-driven papers: **Letta**, born from MemGPT, positions itself as a platform for stateful agents; **Graphiti** focuses on temporally aware context graphs; and **Vertex AI Memory Bank** provides managed cross-session long-term memory in Agent Engine / ADK. These systems matter for deployment and will likely influence future research baselines, even when they are not themselves primary academic papers. [58]

Hermes and OpenClaw

Hermes

Hermes Agent, built by Nous Research [59], is best understood as a frontier **agent runtime with a pragmatic, memory-first systems design**, rather than as a benchmark paper. Its core memory is deliberately bounded and curated: two files, `MEMORY.md` and `USER.md`, are stored under `~/.hermes/`

`memories/`, injected as a frozen snapshot into the system prompt at session start, and managed by the agent through a memory tool. The default documented caps are about 2,200 characters for `MEMORY.md` and 1,375 for `USER.md`; when full, the agent consolidates or replaces entries. Beyond the bounded prompt memory, Hermes stores all sessions in SQLite and uses FTS5-backed `session_search` for unbounded recall across prior CLI and messaging sessions. It can also attach one of eight external memory providers that prefetch relevant memories before turns, sync turns after responses, and mirror built-in writes into the provider. [60]

Architecturally, Hermes combines four memory ideas that the research literature often keeps separate: **working/profile memory** in the system prompt, **episodic recall** via session search, **semantic/profile enrichment** via external providers, and **procedural self-improvement** through a built-in learning loop that creates and improves skills over time. That makes Hermes especially interesting as a prototype of how research concepts migrate into an end-user agent stack. In research terms, Hermes looks closest to a hybrid of bounded semantic/profile memory, episodic search memory, and procedural skill growth. [61]

Its strengths are clear. The bounded memory makes cost predictable. The SQLite/FTS5 session search yields a simple and inspectable episodic substrate. The provider layer lets Hermes inherit newer ideas—knowledge graphs, semantic recall, extracted facts, user modeling—without changing the base agent. And its release cadence indicates active iteration, with an official v0.11.0 release published on April 23, 2026. [62]

Its weaknesses are equally important. The built-in in-prompt memory is intentionally tiny, so performance on complex long-horizon tasks will hinge on the quality of `session_search` and whichever external provider is active. The snapshot model also means the prompt-side built-in memory is not a fully dynamic, continuously re-ranked store. Most importantly for this report's scope, I did not locate a peer-reviewed Hermes paper or a standardized benchmark evaluation in the primary materials reviewed here; public evidence is strongest on architecture and product surface, weaker on controlled experimental claims. [63]

OpenClaw

OpenClaw is also better read as a frontier open-source system than as a canonical academic method. Its memory design is notably **file-native and human-auditable**. Public docs state that the agent remembers by writing plain Markdown files in the workspace: `MEMORY.md` for long-term durable facts, `memory/YYYY-MM-DD.md` daily notes for running context, and optional `DREAMS.md` for reviewable consolidation outputs. Memory tools include `memory_search` and `memory_get`; with embeddings configured, `memory_search` uses hybrid semantic + keyword retrieval. OpenClaw also adds "commitments" as opt-in short-lived future reminders, a memory wiki layer that compiles structured claims/evidence/freshness tracking, an automatic memory flush before compaction, and an optional "dreaming" process that promotes only threshold-qualified items into long-term memory. [64]

This design is unusually aligned with several current research desiderata. It preserves **provenance and inspectability** because memory lives in files. It supports **working memory and recent episodic context** via daily notes, **durable semantic memory** via `MEMORY.md`, **retrieval memory** via hybrid search, and **memory consolidation** via the Dreaming sweep and wiki compilation. Conceptually, OpenClaw sits close to recent work that argues for append-only evidence, retrieval-time synthesis, and clear human oversight rather than opaque continual extraction. It is especially attractive for software-engineering and desktop-assistant settings where users want to inspect or edit memory directly. [65]

Its main strengths are transparency, local-first deployment, composability through memory plugins/backends, and the explicit coupling of compaction and promotion with durable storage. The most obvious weaknesses are that file-backed memory can become brittle without careful maintenance policies, and that fast-moving systems surface versioning risk. In April 2026, a public issue documented a mismatch between Dreaming-related surfaces in public docs and shipped CLI/docs for certain 2026.4.x releases, which is exactly the kind of systems drift that matters for reproducibility. The release stream is active, but users reproducing results should pin exact versions rather than rely on "latest." [66]

In short, Hermes emphasizes **bounded in-prompt memory plus search and pluggable providers**, whereas OpenClaw emphasizes **workspace-native persistence, inspectability, and consolidation hooks**. Both are highly relevant to frontier practice, but neither should be confused with a controlled research benchmark result unless paired with an external evaluation harness such as LoCoMo or MemoryAgentBench. [67]

Benchmarks, Datasets, and Evaluation

The benchmark story is now rich enough to stratify by memory capability.

LoCoMo remains the canonical long-horizon conversational benchmark. Its conversations average roughly 300 turns and 9K tokens over up to 35 sessions, and it supports question answering, event summarization, and multimodal dialogue generation. It is still the closest thing to a shared "ImageNet" for conversational agent memory, which is why so many 2025–2026 papers report LoCoMo. **LongMemEval** complements it with 500 carefully curated questions over scalable user-assistant chat histories and explicitly targets information extraction, multi-session reasoning, temporal reasoning, knowledge updates, and abstention. [68]

The newer generation of benchmarks expands the space in more diagnostic ways. **MemoryAgentBench** reframes evaluation as four competencies—accurate retrieval, test-time learning, long-range understanding, and selective forgetting—and converts long-context data into an incremental multi-turn format more faithful to memory agents. **MemoryArena** then goes one level farther by coupling memory with action in explicitly interdependent multi-session tasks spanning web navigation, planning, search, and formal reasoning. **LifelongAgentBench** targets continual adaptation across database, OS, and knowledge-graph environments. Together, these benchmarks reflect a field-wide realization that memory only matters insofar as it changes future behavior. [69]

Multimodal evaluation is finally becoming serious. **Mem-Gallery** introduces multi-session conversations grounded jointly in text and images and evaluates three functional dimensions: memory extraction/test-time adaptation, memory reasoning, and memory knowledge management. **MIRIX** in turn adds ScreenshotVQA as a high-scale screen-memory setting with nearly 20,000 screenshots per sequence. This matters because many deployed agents increasingly operate over GUIs, screenshots, files, and desktop context, not just chat turns. [70]

The benchmarking literature is also exposing weaknesses in common metrics. Many papers still report **token-level F1**, **BLEU-1**, and **LLM-as-a-Judge** on LoCoMo-style QA. Those are useful, but they are not interchangeable. Memory-R1 explicitly uses F1, BLEU-1, and judge metrics, while LoCoMo-Plus argues that string-matching metrics and explicit task-type prompting can miss more realistic "cue–trigger semantic disconnect” cases where the relevant memory is latent rather than directly queried as a fact. This is one reason cross-paper comparison is delicate. [71]

A practical implication follows. When reading reported state of the art on conversational memory, one should ask at least six questions before trusting the number: which backbone was used, which benchmark variant and subset, whether an adversarial subset was excluded, how memory was chunked/extracted, which judge model scored the outputs, and whether latency/token costs were normalized. Memory-R1, for example, states that it trains on LoCoMo and excludes the adversarial subset in the split it follows; Mem0 emphasizes LLM-as-a-Judge and efficiency; MemMachine emphasizes retrieval-stage optimizations and token reductions; and LoCoMo-Plus explicitly questions whether the standard scoring setup captures cognitive memory at all. [72]

Performance Synthesis, Limitations, and Open Questions

A clear empirical pattern has emerged: **better memory structure and control usually beats simply throwing more context at the model**. The evidence comes from different angles. *MemoryOS* shows sizable LoCoMo gains from hierarchical short-/mid-/long-term storage. *Mem0* shows that a curated persistent memory layer can outperform a full-context baseline while drastically lowering latency and token cost. *LightMem* shows that carefully separating online retrieval from offline consolidation can improve both accuracy and efficiency. *MemMachine* suggests that preserving raw episodic ground truth and moving intelligence to retrieval can outperform more extraction-heavy alternatives. [73]

Another strong pattern is that **time and contradiction handling are now first-class design problems**. The literature has steadily moved away from storing “facts” as static entries. *Temporal Semantic Memory* argues that many systems index dialogue time rather than event time and therefore mishandle durative states. *APEX-MEM* uses append-only storage plus retrieval-time resolution to preserve historical evolution. *G-Memory* and *MIRIX* show that typed, layered, or graph memories are especially valuable when multiple agents or multiple modalities are involved. In other words, the frontier is increasingly about remembering **how the world changed**, not merely what it once contained. [74]

The most exciting methodological shift is **learned memory control**. *Memory-R1* demonstrates that RL can directly optimize memory operations and answer-time distillation, with substantial gains over heuristic baselines and transfer across benchmarks. *Learn to Memorize* and *Mem-α* push in the same direction by treating retrieval, aggregation, or memory construction as trainable policies rather than fixed human-designed rules. If this line continues to scale, the field may move from “memory architectures” to “memory policies” as the dominant unit of innovation. [75]

Important limitations remain. First, many results are still **self-reported** and often compare systems under different base models or evaluation setups. Second, benchmarks remain fragmented: LoCoMo is not MemoryArena, and neither captures the full multimodal or security-governance problem. Third, many systems still depend on lossy LLM extraction at ingest time, which risks compounding mistakes across sessions. Fourth, memory security and governance lag behind capability work; a 2026 survey characterizes attack surfaces across write, store, retrieve, execute, share, and forget/rollback phases and argues that these remain underdeveloped in long-term-memory agents. [76]

The highest-value open questions are these. How should agents decide **what not to remember**? How should they support **rollback, provenance, and deletion** when facts evolve? Can one unify **episodic**

fidelity with **semantic abstraction** without losing traceability? How should memory generalize across **multi-agent coordination**, **GUI/screen interaction**, and **continual learning**? And what is the right trade-off between **human-readable memory substrates** such as files/wiki pages and **machine-optimized latent or graph memories**? Current work has compelling partial answers, but not a consensus solution. [77]

Recommended Experiments and Reproducibility Notes

The next experiments worth doing are no longer generic benchmark reruns. The most useful agenda is a **controlled, same-backbone memory bakeoff**. Choose one open model family and one embedding stack, then evaluate *A-MEM*, *MemoryOS*, *Mem0*, *LightMem*, *MIRIX*, *G-Memory*, and *MemMachine* under a single harness across LoCoMo, LongMemEval, MemoryAgentBench, MemoryArena, and one multimodal task. Keep the backbone, tokenizer, embedding model, retrieval budget, and judge fixed. Without this, many of today's "SOTA" claims remain informative but not decisive. [78]

A second high-payoff experiment is a **write-policy ablation**. Compare heuristic write rules, prompted self-write, supervised write labeling, and RL-trained write/update/delete policies on the same memory substrate. This would directly test whether the recent RL results are mostly due to better retrieval, better write decisions, or better answer-time distillation. Memory-R1, Learn to Memorize, and Mem-α strongly suggest that learned control will matter; what remains unclear is which operation benefits the most from learning. [75]

A third essential experiment is a **temporal contradiction benchmark**. Mix fact updates, durable states, conflicting observations, delayed clarifications, and user preference drift. Then compare graph-temporal systems such as APEX-MEM and Graphiti-like context graphs against summary-first systems and raw episodic systems. The recent temporal-semantic papers imply that this is where many memory agents still fail. [79]

A fourth experiment should target **multimodal memory under storage pressure**. Use ScreenshotVQA or Mem-Gallery and compare raw screenshot retention, visual summarization, typed multimodal memory, and retrieval-only baselines under explicit storage budgets. MIRIX's headline result—a 35% accuracy gain over RAG with 99.9% lower storage—is precisely the kind of claim that deserves independent replication. [80]

For reproducibility, five practices matter more than usual. First, pin the **exact agent/system version** and documentation commit. Hermes has formal published releases, including v0.11.0 on April 23, 2026; MIRIX's v0.1.6 release explicitly repositioned it as a standalone memory API; OpenClaw's fast-moving 2026.4.x line showed documentation/shipping drift around Dreaming surfaces. Second, report the **judge model** separately from the answer model. Third, publish **memory traces**: what was written, updated, deleted, retrieved, and consolidated. Fourth, disclose benchmark choices such as whether adversarial subsets were included. Fifth, report **latency and token cost** alongside accuracy, because many memory systems win precisely by changing the cost-quality frontier rather than only the final score. [81]

The bottom line for practitioners is straightforward. If the goal is a deployable stateful assistant today, prefer a transparent, typed, and benchmarkable external-memory stack with explicit retrieval and consolidation logs. If the goal is frontier research, the most promising work is now at the intersection of **temporal structure**, **multimodal memory**, and **learned memory control**. And if the goal is scientific progress rather than leaderboard churn, the field most urgently needs **shared evaluation harnesses under matched conditions** and **public traces of memory operations**, not just final answer scores. [82]

1 8 9 10 13 14 https://arxiv.org/abs/2304.03442
https://arxiv.org/abs/2304.03442

2 30 31 78 https://arxiv.org/abs/2502.12110
https://arxiv.org/abs/2502.12110

3 68 https://arxiv.org/abs/2402.17753
https://arxiv.org/abs/2402.17753

4 42 43 80 https://arxiv.org/abs/2507.07957
https://arxiv.org/abs/2507.07957

5 15 https://github.com/joonspk-research/generative_agents
https://github.com/joonspk-research/generative_agents

6 35 https://github.com/BAI-LAB/MemoryOS
https://github.com/BAI-LAB/MemoryOS

7 16 17 https://arxiv.org/abs/2305.16291
https://arxiv.org/abs/2305.16291

11 24 25 https://arxiv.org/abs/2310.08560
https://arxiv.org/abs/2310.08560

12 76 https://arxiv.org/abs/2602.10715
https://arxiv.org/abs/2602.10715

18 57 https://github.com/minedojo/voyager
https://github.com/minedojo/voyager

19 20 https://arxiv.org/abs/2303.11366
https://arxiv.org/abs/2303.11366

21 https://github.com/noahshinn/reflexion
https://github.com/noahshinn/reflexion

22 https://arxiv.org/abs/2305.10250
https://arxiv.org/abs/2305.10250

23 https://github.com/zhongwanjun/MemoryBank-SiliconFriend
https://github.com/zhongwanjun/MemoryBank-SiliconFriend

26 https://github.com/deductive-ai/MemGPT
https://github.com/deductive-ai/MemGPT

27 28 59 https://arxiv.org/abs/2308.10144
https://arxiv.org/abs/2308.10144

29 https://github.com/LeapLabTHU/ExpeL
https://github.com/LeapLabTHU/ExpeL

32 https://github.com/agiresearch/a-mem
https://github.com/agiresearch/a-mem

33 34 73 https://arxiv.org/abs/2506.06326
https://arxiv.org/abs/2506.06326

[36] [37] https://arxiv.org/abs/2504.19413
https://arxiv.org/abs/2504.19413

[38] https://github.com/mem0ai/mem0
https://github.com/mem0ai/mem0

[39] [40] https://arxiv.org/abs/2506.07398
https://arxiv.org/abs/2506.07398

[41] https://github.com/bingreeky/GMemory
https://github.com/bingreeky/GMemory

[44] https://github.com/Mirix-AI/MIRIX
https://github.com/Mirix-AI/MIRIX

[45] [46] [74] https://arxiv.org/abs/2601.07468
https://arxiv.org/abs/2601.07468

[47] [48] https://arxiv.org/abs/2604.07798
https://arxiv.org/abs/2604.07798

[49] https://github.com/zjunlp/LightMem
https://github.com/zjunlp/LightMem

[50] [51] [71] [72] [75] https://arxiv.org/html/2508.19828v5
https://arxiv.org/html/2508.19828v5

[52] [53] [79] https://arxiv.org/abs/2604.14362
https://arxiv.org/abs/2604.14362

[54] [55] https://arxiv.org/abs/2604.04853
https://arxiv.org/abs/2604.04853

[56] https://github.com/MemMachine/MemMachine
https://github.com/MemMachine/MemMachine

[58] https://www.letta.com/about-us
https://www.letta.com/about-us

[60] [62] [63] [67] https://github.com/NousResearch/hermes-agent/blob/main/website/docs/user-guide/features/memory.md
https://github.com/NousResearch/hermes-agent/blob/main/website/docs/user-guide/features/memory.md

[61] https://github.com/nousresearch/hermes-agent
https://github.com/nousresearch/hermes-agent

[64] [65] [66] https://github.com/openclaw/openclaw/blob/main/docs/concepts/memory.md
https://github.com/openclaw/openclaw/blob/main/docs/concepts/memory.md

[69] [77] [82] https://arxiv.org/abs/2507.05257
https://arxiv.org/abs/2507.05257

[70] https://arxiv.org/abs/2601.03515
https://arxiv.org/abs/2601.03515

[81] https://github.com/NousResearch/hermes-agent/releases
https://github.com/NousResearch/hermes-agent/releases